

# BK-TREE: EFFICIENT RETRIEVAL OF SIMILAR STRINGS

---

*Session 5*

*David Buchaca Prats*  
*2023*

# DEALING WITH WORDS OUT OF THE VOCABULARY

---

- Based on two notions that we have already seen: The edit distance and a Language model, we can build a spellchecker.
- Let us consider that we want to correct misspelled words:
  - That is: words that are not in the vocabulary
- Base algorithm: Let  $x$  be a sentence and  $V$  the vocabulary.

For  $w$  in  $x$ :

    If  $w$  not in  $V$ :

        Find the closest words to  $w$  (candidate search)

        Evaluate each of the candidate words (candidate evaluation)

        Return the most probable candidate

# FINDING THE CLOSEST ITEMS TO A QUERY

---

- This is an extremely relevant problem for many Data Science and ML problems.
- Sklearn implements the kdtree for dense vectors.
- What if we have strings...?

```
n_examples = 3_000_000
n_features = 25

X,y = sklearn.datasets.make_blobs(n_examples, n_features, centers=30, random_state=123)
x = X[0:1]
```

```
kdtree = sklearn.neighbors.KDTree(X)
kdtree
```

```
<sklearn.neighbors._kd_tree.KDTree at 0x7fa1778e9810>
```

```
%timeit kdtree.query(x, return_distance=False)
```

```
33.1 µs ± 220 ns per loop (mean ± std. dev. of 7 runs, 10,000 loops each)
```

```
%timeit closest_match = np.argpartition(np.sum((X - x)**2,axis=1),1)[0]
```

```
301 ms ± 9.49 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

```
kdtree.query(x, return_distance=False)
```

```
array([[0]])
```

```
np.argpartition(np.sum((X - x)**2,axis=1),1)[0]
```

# FINDING THE CLOSEST ITEMS TO A QUERY

---

- Efficient search of similar values in a dataset is a very challenging problem. In particular, computing distances between a word and a huge vocabulary can be computationally expensive.
- Let  $w$  be a string that is out of the vocabulary.
- Let us consider  $W_k(w; X) = \{w_j \mid w_j \in V, d(w, w_j) < k\}$
- Finding  $W_k(w; X)$  can be done in two different ways:
  - I) compute  $d(w, w_j)$  for all  $w_j$  then select the elements that are at distance at most  $k$
  - II) Use a data structure to avoid computing  $d(w, w_j)$  for all  $w_j$  in  $X$ .

# TREE INTUITION

---

- We can build a tree to do efficient search of similar words. This will allow us to prune a lot of the search space, with the objective of avoiding many distance computations on a big part of the vocabulary.
- Example: consider  $w = \text{pleistation}$ 
  - ana                       $d(\text{pleistation}, \text{ana})$
  - playstation           $d(\text{pleistation}, \text{playstation})$
  - house                   $d(\text{pleistation}, \text{house})$
- If pleistation has 11 characters and we want all candidates to be at most at distance  $k=3$ , is there any need to compute  $d(\text{pleistation}, \text{ana})$  ?
  - Ana has 3 characters!

# BK-TREE

---

- To create a BK-tree we will follow this approach.
  - Select any word from the vocabulary and use it as the root node.
  - Keep adding words until all vocabulary is in the tree.
    - Each time we add a word the distance between the word and the root node is computed, let us assume this distance is  $d$ .
    - If no node from the root node is at distance  $d$  we add a new leaf as a descendant of the root node with edge value equal to  $d$ .
    - If there exists another node at distance  $d$  then... we repeat this process redefining the root node as the node that produced the collision (pleistation, ana)




# BK TREE CONSTRUCTION EXAMPLE

---


- Let us consider the data [book, books, cake], we start from **book** (which becomes root node)



- [book, books, cake]: **books** comes and we compute  $d(\text{book}, \text{books}) = 1$




- [book, books, cake]: **cake** comes and we compute  $d(\text{book}, \text{cake}) = 4$



# BK-TREE EXAMPLE: [BOOK, BOOKS, CAKE, BOO, CAPE, CART, BOON, COOK]

---

- [book, books, cake, boo]: **boo** comes and we compute  $d(\text{book}, \text{boo}) = 1$ .  
Note that there is already **books** at distance 1.
- The BK tree has to respect that every node have all children with different distances, **since there is already a word at the same edit distance 1** we go to the branch of words at distance 1.
- If there is a **collision** (like we have now) the new word must become a children of the collided word. In this case, a children of book.
- The new weight from books to boo has to be  $\text{distance}(\text{books}, \text{boo}) = 2$ .






# BK-TREE EXAMPLE: [BOOK, BOOKS, CAKE, BOO, CAPE, CART, BOON, COOK]

---


- Root=book: Compute  $d(\text{book}, \text{cape}) = 4$ 
  - Collision! There is already cake at distance 4 from book
  - Root node is now cake
  - Root=cake: Compute  $d(\text{cake}, \text{cape}) = 1$
  - There is no descendant from cake at distance 1  $\Rightarrow$  we can add it



# BK-TREE EXAMPLE: [BOOK, BOOKS, CAKE, BOO, CAPE, CART, BOON, COOK]

---

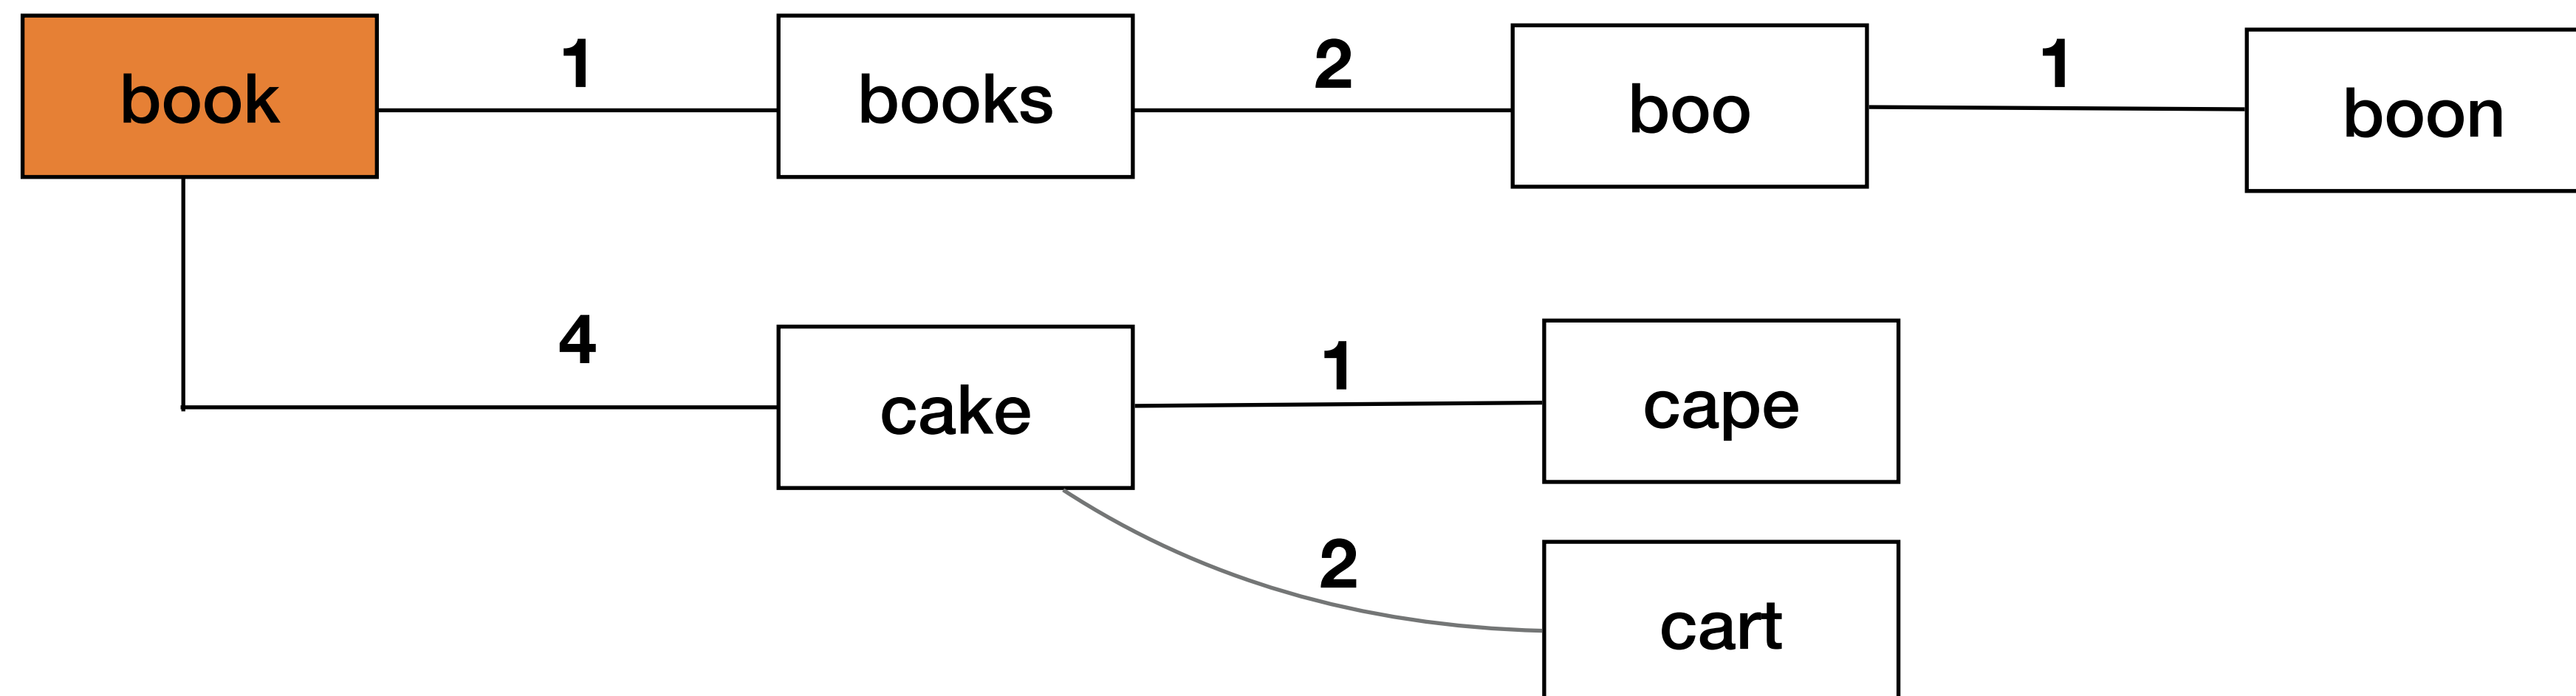
- Root=book: Compute  $d(\text{book}, \text{cart}) = 4$ 
  - Collision! There is already cake at distance 4 from book
  - Root node is now cake
  - Root=cake: Compute  $d(\text{cake}, \text{cart}) = 2$
  - There is no descendant from cake at distance 2  $\Rightarrow$  we can add it



# BK-TREE EXAMPLE: [BOOK, BOOKS, CAKE, BOO, CAPE, CART, BOON, COOK]

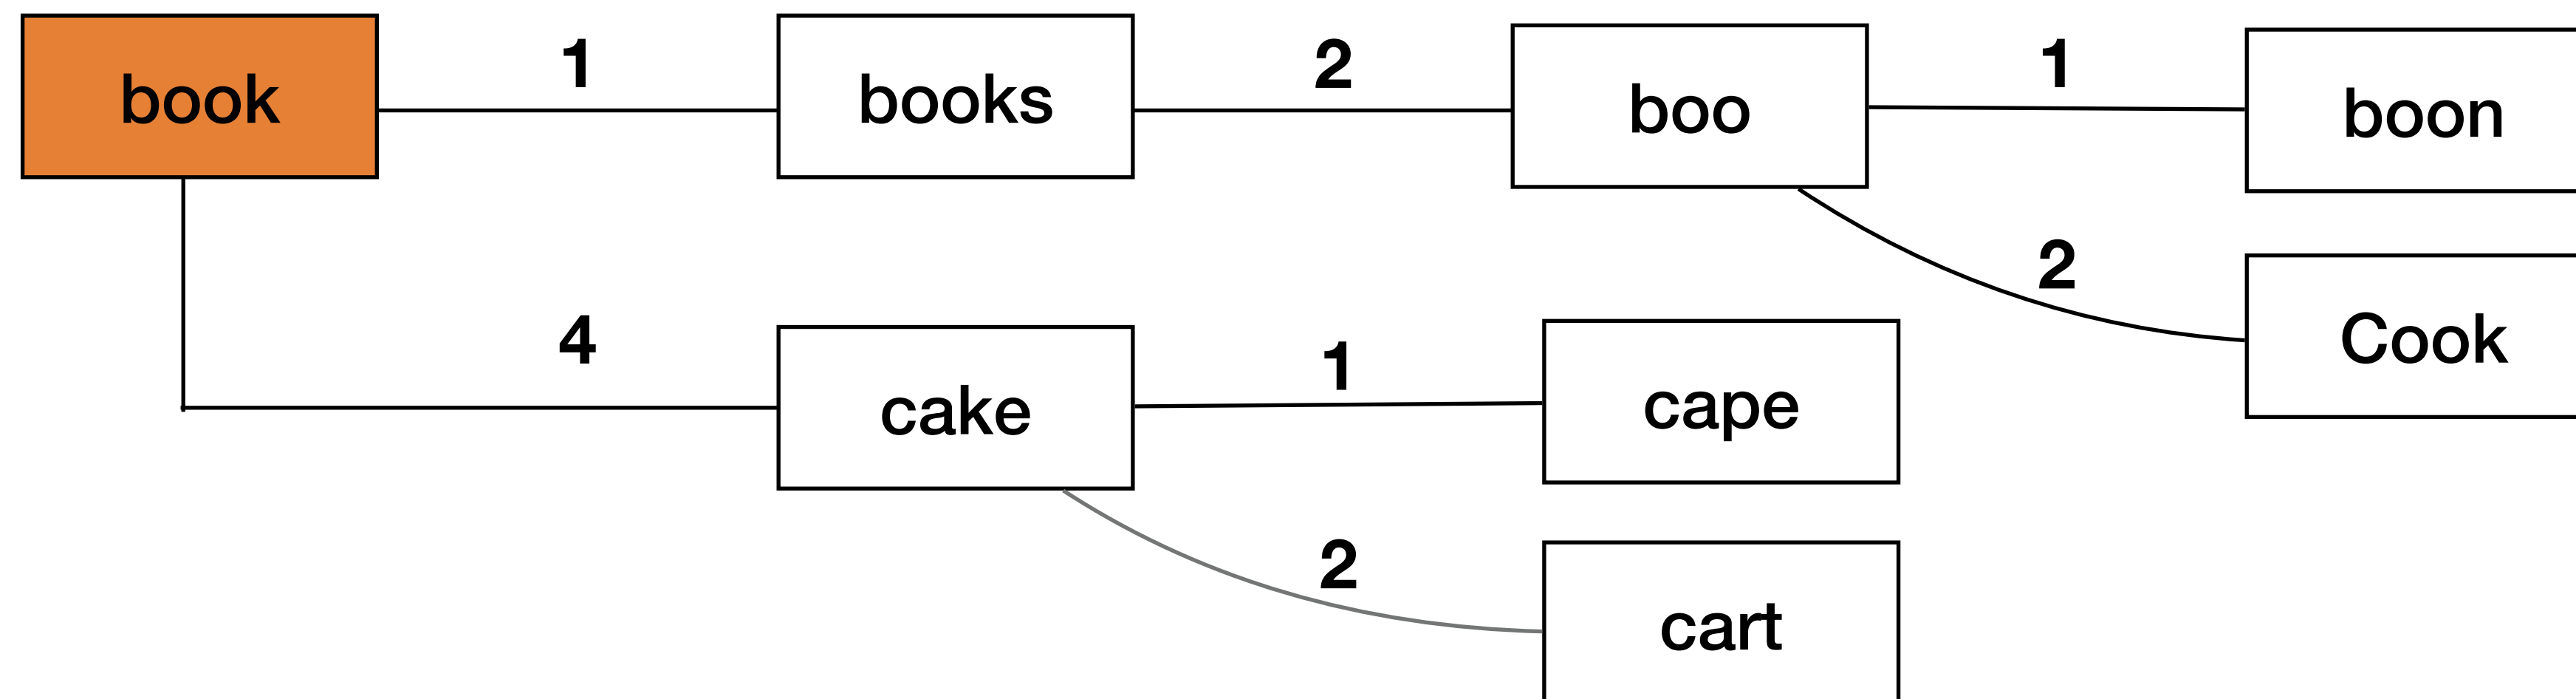
---

- Root=book: Compute  $d(\text{book}, \text{boon}) = 1$ 
  - Collision! There is already books at distance 1 from book
  - Root node is now books
  - Root=books: Compute  $d(\text{books}, \text{boon}) = 2$
  - Collision! There is already boo at distance 2 from books
  - Root=boo: Compute  $d(\text{books}, \text{boon}) = 1$
  - There is no descendant from boo at distance 1  $\Rightarrow$  we can add it



# BK-TREE EXAMPLE: [BOOK, BOOKS, CAKE, BOO, CAPE, CART, BOON, COOK]


- Root=book: Compute  $d(\text{book}, \text{cook}) = 1$ 
  - Collision! There is already books at distance 1 from book
  - Root node is now books
  - Root=books: Compute  $d(\text{books}, \text{cook}) = 2$
  - Collision! There is already boo at distance 2 from books
  - Root=boo: Compute  $d(\text{boo}, \text{cook}) = 2$
  - There is no descendant from boo at distance 2 => we can add it



# BK TREE: STORAGE IN MEMORY

---


- Let us consider the following tree
- What would be a reasonable way to store the tree in memory?



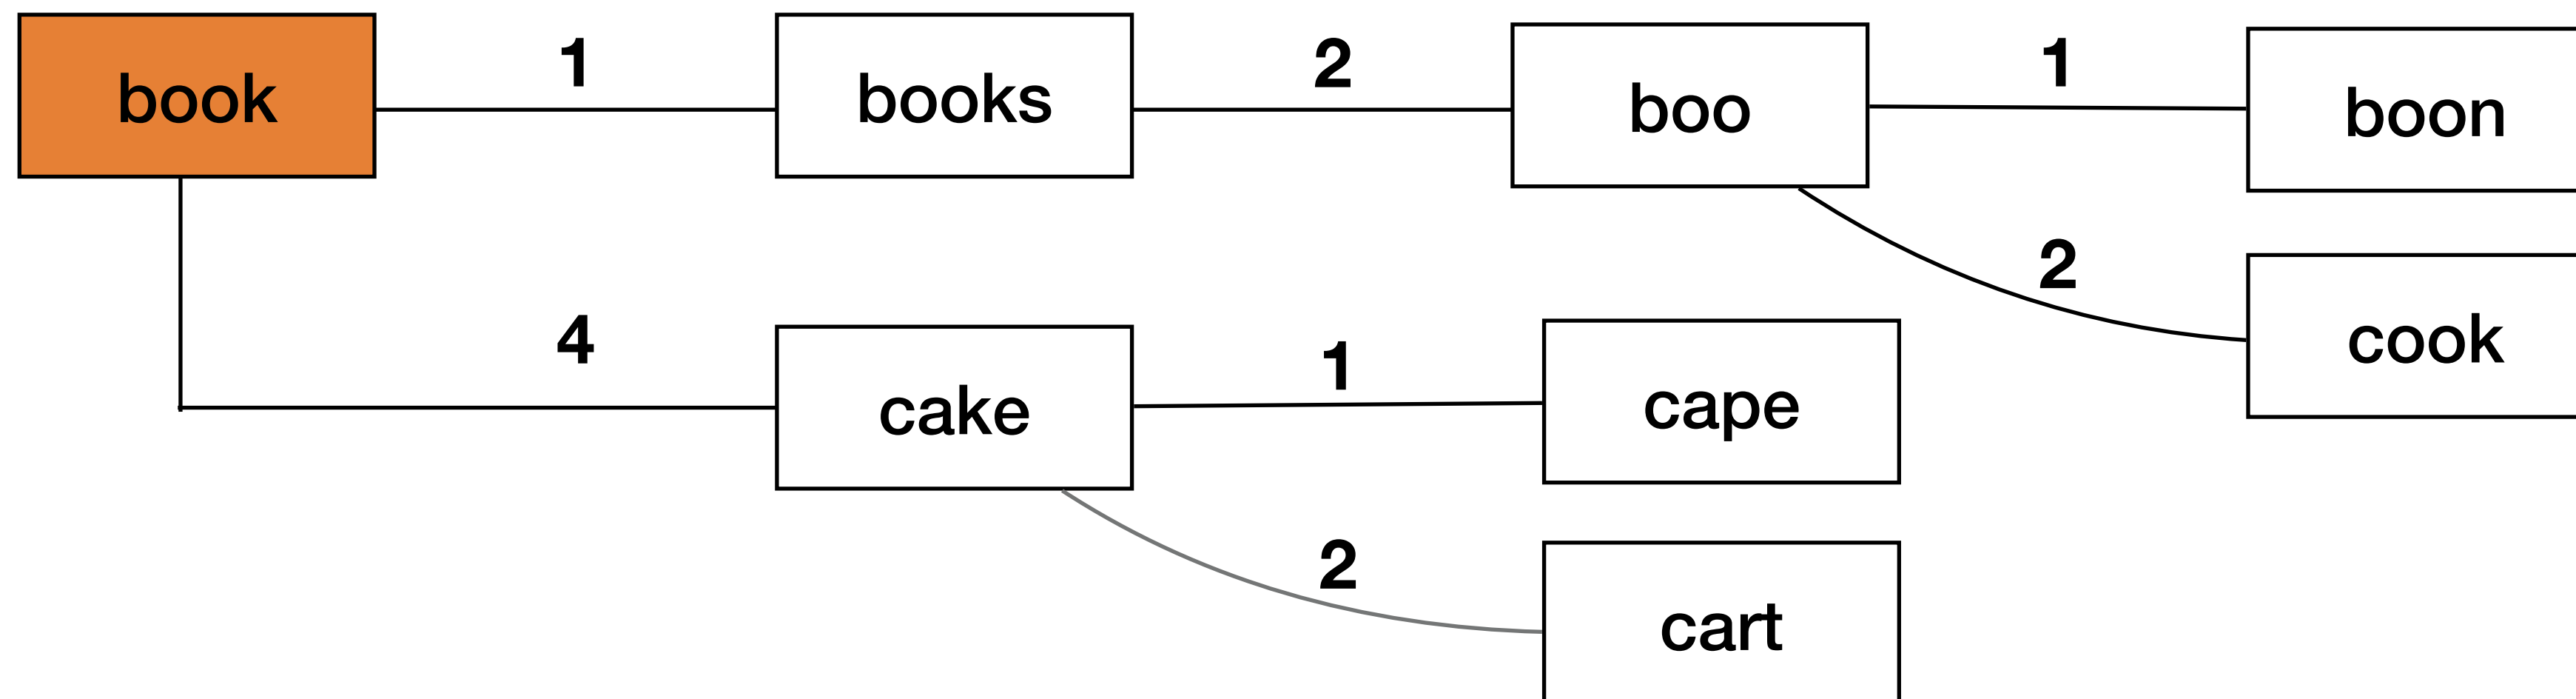
- A sensible way could be a tuple
- The first element is the word assigned to the node
- The second element is the subtree that spawns from that node
  - A subtree can be represented as a `Dict[Int, Tuple]`
    - keys are the distances to the root node
    - Values are tuples which represent subtree
- In other words: `('book', {1: ('books', {}), 4: ('cake', {})})`

# BK TREE: STORAGE IN MEMORY

---



- ('book',  
{1: ('books', {2: ('boo', {})}),  
4: ('cake', {})}))



- ('book',  
{1: ('books', {2: ('boo', {1: ('boon', {}), 2: ('cook', {})}),  
4: ('cake', {1: ('cape', {}), 2: ('cart', {})}))})



# SEARCHING IN A BK-TREE

- Problem: Search all words that appear at distance less or equal than a tolerance  $T$  from a query word  $q$ .
- Bad solution: Compute all edit distances between  $q$  and  $w$  for  $w$  in the Vocabulary.
- Key idea: Visit all words  $w$  that are at distance  $[d(w,q)-T, d(w,q)+T]$ .

- Example:

- $q = \text{vook}, T = 2$

- $d(\text{vook}, \text{book}) = 1$

- Consider  $w$  from key 4 from book (yellow)

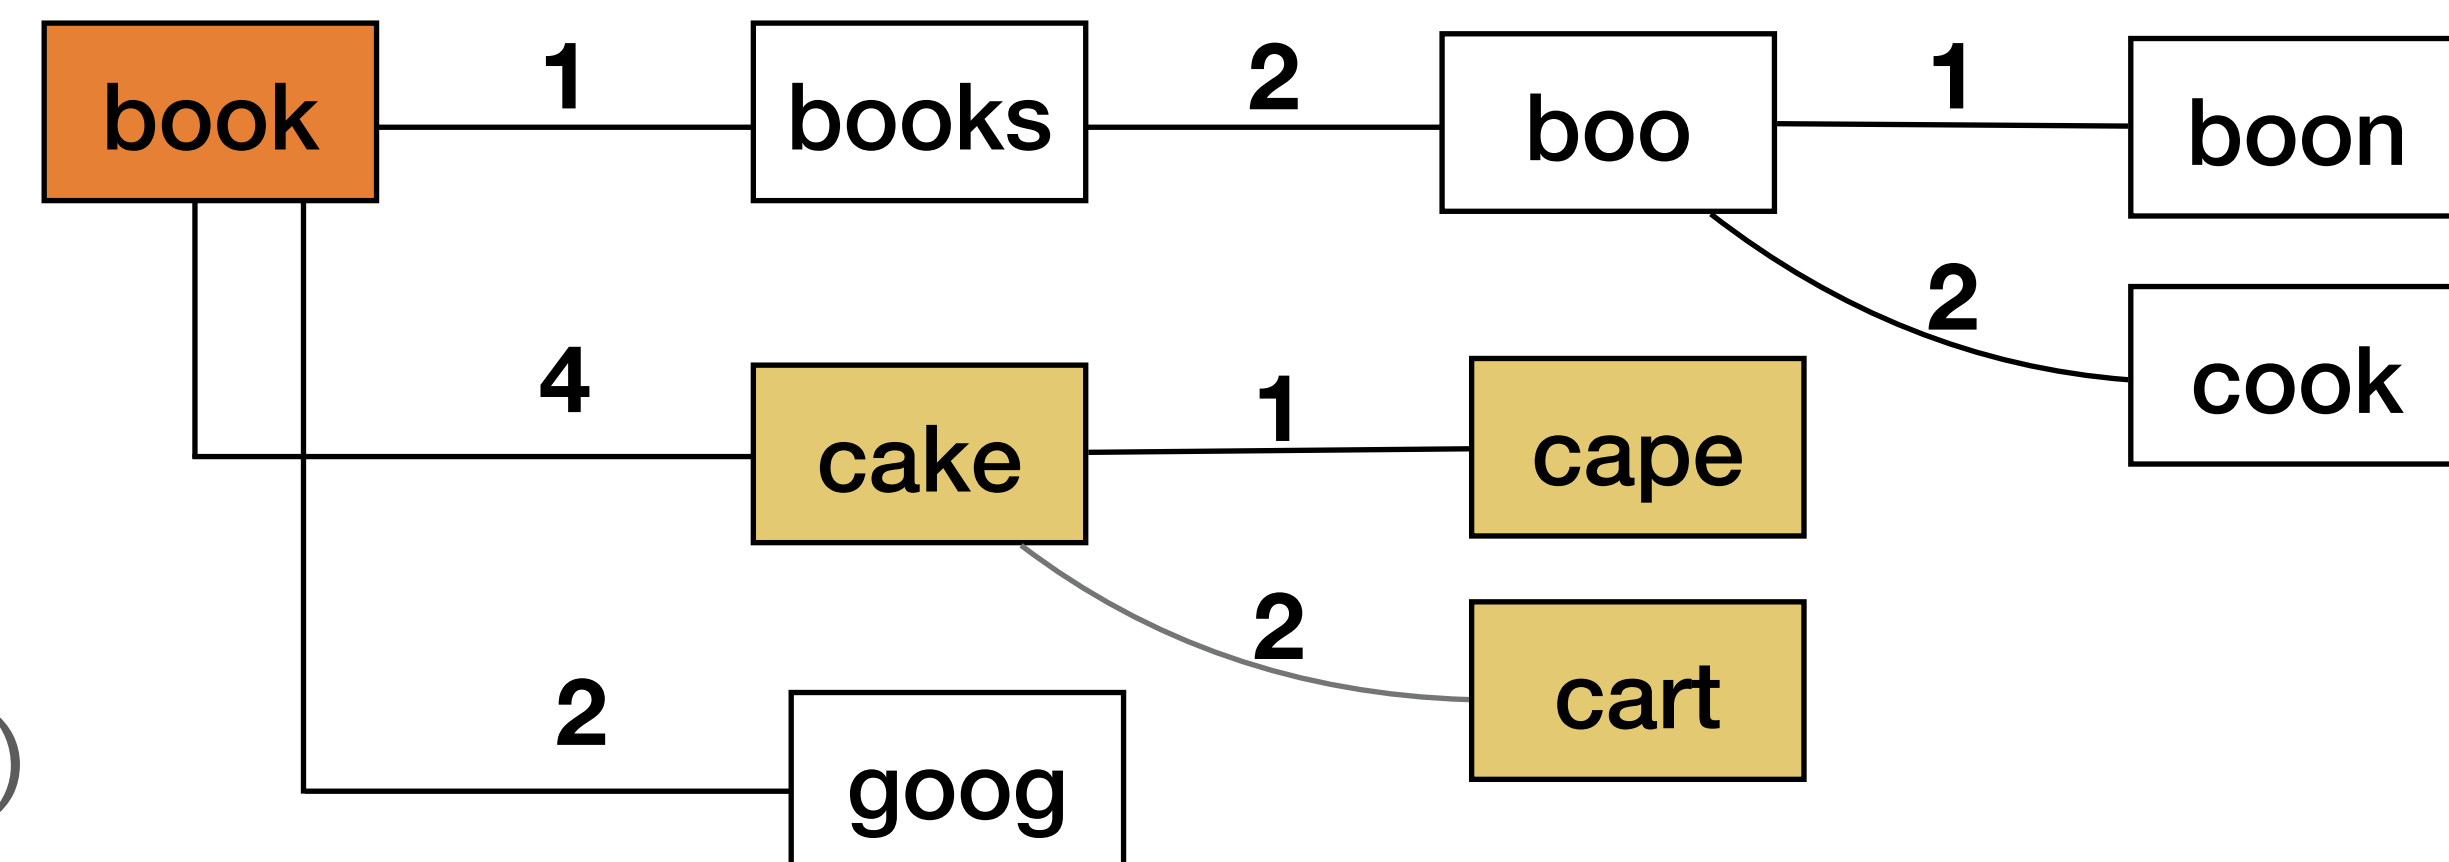
- By construction all words in yellow subtree are at a distance 4 from book

- $d(\text{vook}, w) \leq d(\text{vook}, \text{book}) + d(\text{book}, w) = 5$

This is 1

This is 4

No need to search for  $w$  in the yellow subtree: we want  $d(\text{vook}, w) \leq 2$



# BK-TREE EXAMPLE: SEARCHING

➤ Let us consider

➤  $q = \text{cage}$ ,  $T = 1$ , candidates = []


➤ Select candidate **book** from search=[book]

➤  $d(\text{book}, \text{cage}) = 4 \Rightarrow$  candidates is not updated

➤ Crawl all children of book at distance  $I = [4-1, 4+1] = [3, 5]$


➤ Only node cake is connected to book and with distance in  $I = [3, 5]$

➤ search = [book, cake] \ book = [cake]



# BK-TREE EXAMPLE: SEARCHING

- Let us consider
- $q = \text{cage}$ ,  $T = 1$ , candidates = []
- Select candidate **cake** from search=[cake]
  - $d(\text{cake}, \text{cage}) = 1 \Rightarrow \text{candidates} += [\text{cake}]$ 
    - Crawl all children of cake at distance  $I = [1-1, 1+1] = [0, 2]$
    - Only are 2 possible nodes, search=[cape, cart]



# BK-TREE EXAMPLE: SEARCHING

➤ Let us consider

➤  $q = \text{cage}$ ,  $T = 1$ , candidates = [cake]


➤ Select candidate **cape** from search=[cape, cart]

➤  $d(\text{cape}, \text{cage}) = 1 \Rightarrow \text{candidates} += [\text{cape}]$

➤ Crawl all children of cape at distance  $I = [1-1, 1+1] = [0, 2]$

➤ cape has no children

➤ search = [cape, cart] \ cape = [cart]



# BK-TREE EXAMPLE: SEARCHING

➤ Let us consider

➤  $q = \text{cage}$ ,  $T = 1$ , candidates = [cake, cape]

➤ Select candidate **cart** from search=[cart]


➤  $d(\text{cart}, \text{cage}) = 2 \Rightarrow$  candidates is not updated

➤ Crawl all children of cape at distance  $I = [1-1, 1+1] = [0, 2]$

➤ Cage has no children


➤ search = [cart] \ cart = []  $\Rightarrow$  Search space is empty, stop search

➤ The resulting set of possible candidates at distance 1 are: [cape, cake]



# BK-TREE EXAMPLE: SEARCHING

- To sum up:
  - We started from:
    - $q = \text{cage}$ ,  $T = 1$ ,  $\text{candidates} = []$ ,  $\text{search} = [\text{book}]$
  - After searching in the BK-Tree we know
    - The set of possible candidates at distance 1 are:  $[\text{cape}, \text{cake}]$ .
- Observation: we ended up computing 4 edit distances yet we have 8 nodes.





# BK- TREE SPEEDUP

---

- In the case that the search space is drastically pruned, the speedup can be massive:

```
word = "anthropomorphologically"  
max_dist = 2  
sort_candidates=False
```

```
%timeit candidates_ext = get_candidates_exhaustive(word,max_dist,words)
```

404 ms ± 5.26 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
candidates_ext = get_candidates_exhaustive(word,max_dist,words)  
candidates_ext
```

```
[(1, 'anthropomorphological'), (1, 'anthropomorphologically')]
```

```
word = "anthropomorphologically"
```

```
%timeit candidates_ext = t.query(word, 2)
```

214  $\mu$ s ± 1.77  $\mu$ s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```
candidates_ext = t.query(word, 2)  
candidates_ext
```

```
[(1, 'anthropomorphological'), (1, 'anthropomorphologically')]
```

# BK- TREE SPEEDUP

---

- If the pruned search space still contains a huge amount of words the speedup might not be that huge:

```
word = "astrologi"  
max_dist = 2  
sort_candidates=False  
  
%timeit candidates_ext = get_candidates_exhaustive(word, max_dist, words)
```

319 ms ± 1.8 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
word = "astrologi"  
  
%timeit t.query(word, 2)
```

96.2 ms ± 737 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)