

Hidden Markov Models (II)

Natural Language Processing

David Buchaca Prats, Daniel Ortiz Martínez

Forward Backward Algorithm

- Forward quantity

$$f(i, x, c) := P(X_1 = x_1, \dots, X_i = x_i, Y_i = c)$$

- Backward quantity

$$b(i, x, c) := P(X_{i+1} = x_{i+1}, \dots, X_N = x_N | Y_i = c)$$

Efficient Forward Quantity Computations

- If we define $f(i, x, c) := P(X_1 = x_1, \dots, X_i = x_i, Y_i = c)$ then, the forward quantity at position i can be written as:

$$f(i, x, c) = P_{\text{emiss}}(x_i | c) \cdot \left(\sum_{\tilde{c} \in \Lambda} P_{\text{trans}}(c | \tilde{c}) \cdot f(i-1, x, \tilde{c}) \right)$$

Forward Quantities Definition

- Let us consider a sequence x with N words
- For every state c let us define:

$$f(1, x, c) = P_{\text{init}}(c_k | \text{start}) \cdot P_{\text{emiss}}(x_1 | c_k)$$

- For every position $i = 2$ up to N

$$f(i, x, c) = P_{\text{emiss}}(x_i | c) \cdot \left(\sum_{\tilde{c} \in \Lambda} P_{\text{trans}}(c | \tilde{c}) \cdot f(i-1, x, \tilde{c}) \right)$$

- For position $N + 1$

$$f(N+1, \text{stop}) = \sum_{c_l \in \Lambda} P_{\text{final}}(\text{stop} | c) \cdot f(N, x, c)$$


Forward Quantities

$$f(i, x, c) = P_{\text{emiss}}(x_i | c) \cdot \left(\sum_{\tilde{c} \in \Lambda} P_{\text{trans}}(c | \tilde{c}) \cdot f(i-1, x, \tilde{c}) \right)$$

Λ							
	$f(1, x, \cdot)$	$f(2, x, \cdot)$		$f(i-1, x, \cdot)$	$f(i, x, \cdot)$		$f(N, x, \cdot)$


Forward Quantities

$$f(i, x, c) = P_{\text{emiss}}(x_i | c) \cdot \left(\sum_{\tilde{c} \in \Lambda} P_{\text{trans}}(c | \tilde{c}) \cdot f(i-1, x, \tilde{c}) \right)$$




Backward Quantities

$$b(i, x, c) = \sum_{\tilde{c} \in \Lambda} P_{\text{trans}}(\tilde{c}|c) \cdot b(i+1, x, \tilde{c}) \cdot P_{\text{emiss}}(x_{i+1}|\tilde{c})$$



Backward Quantities

$$b(i, x, c) = \sum_{\tilde{c} \in \Lambda} P_{\text{trans}}(\tilde{c}|c) \cdot b(i+1, x, \tilde{c}) \cdot P_{\text{emiss}}(x_{i+1}|\tilde{c})$$



Forward Backward Algorithm

- Proposition I: $P(X = x, Y_i = c) = f(i, x, c) \cdot b(i, x, c)$
- Proposition II: $P(X = x) = \sum_{c \in \Lambda} f(i, x, c) \cdot b(i, x, c)$
- Proposition III: $P(Y_i = c | X = x) = \frac{f(i, x, c) \cdot b(i, x, c)}{P(X=x)}$
- Proposition IV (Posterior decoding):

$$y^* = \arg \max_{c \in \Lambda} P(Y_i = c | X_1 = x_1, \dots, X_N = x_N) = \arg \max_{c \in \Lambda} f(i, x, c) \cdot b(i, x, c)$$

Proposition 1

- Proposition 1: $P(X = x, Y_i = c) = f(i, x, c) \cdot b(i, x, c)$
- Remember that we defined

$$b(i, x, c) := P(X_{i+1} = x_{i+1}, \dots, X_N = x_N | Y_i = c)$$

$$f(i, x, c) := P(X_1 = x_1, \dots, X_i = x_i, Y_i = c)$$

- Therefore

$$P(X = x, Y_i = c) = P(X_1 = x_1, \dots, X_N = x_N, Y_i = c) =$$

$$P(X_1 = x_1, \dots, X_i = x_i, \dots, X_N = x_N, Y_i = c) =$$

$$P(X_1 = x_1, \dots, X_i = x_i | X_{i+1} = x_{i+1}, \dots, X_N = x_N, Y_i = c) \cdot P(X_{i+1} = x_{i+1}, \dots, X_N = x_N, Y_i = c) =$$

$$P(X_1 = x_1, \dots, X_i = x_i | Y_i = c) \cdot P(X_{i+1} = x_{i+1}, \dots, X_N = x_N, Y_i = c) =$$

$$b(i, x, c) \cdot f(i, x, c)$$

Proposition II

- Proposition II: $P(X = x) = \sum_{c \in \Lambda} f(i, x, c) \cdot b(i, x, c)$

$$P(X = x) = \sum_{c \in \Lambda} P(X = x, Y_i = c) \stackrel{(\text{Prop. I})}{=} \sum_{c \in \Lambda} f(i, x, c) \cdot b(i, x, c)$$

Proposition III

- Proposition III: $P(Y_i = c | X = x) = \frac{f(i, x, c) \cdot b(i, x, c)}{P(X = x)}$

$$P(Y_i = c | X = x) = \frac{P(Y_i = c, X = x)}{P(X = x)} \stackrel{(\text{Prop.I})}{=} \frac{f(i, x, c) \cdot b(i, x, c)}{P(X = x)}$$

Proposition IV

- Proposition IV (Posterior decoding):

$$y^* = \arg \max_{c \in \Lambda} P(Y_i = c | X_1 = x_1, \dots, X_N = x_N) = \arg \max_{c \in \Lambda} f(i, x, c) \cdot b(i, x, c)$$

- Proposition III tells us $P(Y_i = c | X = x) = \frac{f(i, x, c) \cdot b(i, x, c)}{P(X=x)}$
- Therefore

$$\begin{aligned} y^* &= \arg \max_{c \in \Lambda} P(Y_i = c | X_1 = x_1, \dots, X_N = x_N) = \\ &\arg \max_{c \in \Lambda} \frac{f(i, x, c) \cdot b(i, x, c)}{P(X = x)} = \\ &\arg \max_{c \in \Lambda} f(i, x, c) \cdot b(i, x, c) \end{aligned}$$

Forward Backward Algorithm

```
forward_backward( $x, \Lambda, P_{\text{init}}, P_{\text{trans}}, P_{\text{final}}, P_{\text{emiss}}$ )  
  # Forward pass  
  for  $c_k \in \Lambda$ :  
    forward( $1, x, c_k$ ) =  $P_{\text{init}}(c_k | \text{start}) \cdot P_{\text{emiss}}(x_1 | c_k)$   
  for  $i = 1$  to  $N$ :  
    for  $\tilde{c} \in \Lambda$ :  
      forward( $i, x, \tilde{c}$ ) =  $\left( \sum_{c_k \in \Lambda} P_{\text{trans}}(\tilde{c} | c_k) \cdot \text{forward}(i - 1, x, c_k) \right) P_{\text{emiss}}(x_i | \tilde{c})$   
  
  # Backward pass  
  for  $c_k \in \Lambda$ :  
    backward( $N, x, c_k$ ) =  $P_{\text{final}}(\text{stop} | c_k)$   
  for  $i = N - 1$  to  $1$ :  
    for  $\tilde{c} \in \Lambda$ :  
      backward( $i, x, \tilde{c}$ ) =  $\left( \sum_{c_k \in \Lambda} P_{\text{trans}}(c_k | \tilde{c}) \cdot \text{backward}(i + 1, x, c_k) \right) P_{\text{emiss}}(x_{i+1} | c_k)$ 
```

Viterbi Decoding

- So far we have seen how to perform posterior decoding:

$$y_i^* = \arg \max_{y_i \in \Lambda} P(Y_i = y_i | X_1 = x_1, \dots, X_N = x_N)$$

- We define Viterbi decoding as

$$y^* = \arg \max_{y \in \Lambda^N} P(Y_1 = y_1, \dots, Y_N = y_N | X_1 = x_1, \dots, X_N = x_N)$$

- Note that conditional and joint probabilities are proportional

$$\arg \max_{y \in \Lambda^N} P(Y = y | X = x) = \arg \max_{y \in \Lambda^N} P(Y = y, X = x)$$

- Viterbi decoding is very similar to the forward pass of the FB algorithm
- It makes use of the same trellis structure to efficiently represent the exponential number of sequences without prohibitive cost
- The only difference from the FB algorithm is the recursion, that takes the maximum instead of summing over all possible hidden states

$$\text{viterbi}(i, x, c) := \max_{y_1, \dots, y_i} P(Y_1 = y_1, \dots, Y_i = y_i, X_1 = x_1, \dots, X_i = x_i)$$

Viterbi Quantities

- For every state c let us define:

$$\text{viterbi}(1, x, c) = P_{\text{init}}(c|\text{start}) \cdot P_{\text{emiss}}(x_1|c)$$

- For every position $i = 2$ up to $N - 1$:

$$\text{viterbi}(i, x, c) = P_{\text{emiss}}(x_i|c) \cdot \max_{\tilde{c} \in \Lambda} (P_{\text{trans}}(c|\tilde{c}) \cdot \text{viterbi}(i - 1, x, \tilde{c}))$$

- For every state at the last position

$$\text{viterbi}(N, x, c) = P_{\text{emiss}}(x_N|c) \cdot \max_{\tilde{c} \in \Lambda} (P_{\text{trans}}(c|\tilde{c}) \cdot \text{viterbi}(N - 1, x, \tilde{c}))$$

- We can define the Viterbi quantity at the stop position as:

$$\text{viterbi}(N + 1, x, \text{stop}) = \max_{c \in \Lambda} P_{\text{final}}(\text{stop}|c) \cdot \text{viterbi}(N, x, c)$$

- Using the recurrence rule:

$$\max_{y \in \Lambda^N} P(X = x, Y = y) = \max_{c \in \Lambda} P_{\text{final}}(\text{stop}|c) \cdot \text{viterbi}(N, x, c)$$

- The Viterbi algorithm tells us:


$$\text{viterbi}(N + 1, x, \text{stop}) := \max_{y_1, \dots, y_N} P(Y_1 = y_1, \dots, Y_N = y_N, X_1 = x_1, \dots, X_i = x_N)$$

Demonstration Viterbi at Pos. $N+1$ is the Max. Probability

$$\text{viterbi}(N + 1, x, \text{stop}) := \max_{y \in \Lambda^N} P(X = x, Y = y)$$

Exercise!

Viterbi Algorithm Diagram



Viterbi Algorithm Forwards

- Previous recurrence finds the hidden state sequence y^* with highest probability $P(X = x, Y = y^*)$
- To sum up, we have seen that if we consider the quantities:

$$\text{viterbi}(i, x, y_i) = \max_{y_1, \dots, y_{i-1}} P(Y_1 = y_1, \dots, Y_i = y_i, X_1 = x_1, \dots, X_i = x_i)$$

- We can compute $\text{viterbi}(i, x, c)$ using $\text{viterbi}(i - 1, x, \cdot)$

$$\text{viterbi}(i, x, c) = P_{\text{emiss}}(x_i | c) \cdot \max_{\tilde{c} \in \Lambda} (P_{\text{trans}}(c | \tilde{c}) \cdot \text{viterbi}(i - 1, x, \tilde{c}))$$

- Moreover

$$\text{viterbi}(N + 1, x, \text{stop}) = \max_{c_I \in \Lambda} P_{\text{final}}(\text{stop} | c_I) \cdot \text{viterbi}(N, x, c_I) = \max_{c_1, \dots, c_N \in \Lambda} P(X = x, Y = y)$$

Viterbi Algorithm Backwards

- Once the Viterbi value at position N is computed the algorithm can backtrack using the following recurrence:

$$\text{backtrack}(N + 1, x, \text{stop}) = \arg \max_{c_l \in \Lambda} P_{\text{final}}(\text{stop} | c_l) \cdot \text{viterbi}(N, x, c_l)$$

$$\text{backtrack}(i, x, c) = \arg \max_{\tilde{c} \in \Lambda} (P_{\text{trans}}(c | \tilde{c}) \cdot \text{viterbi}(i - 1, x, \tilde{c}))$$

- To do this we need to keep track of the backtrack quantities when we compute the viterbi quantities

Decoding with the Viterbi Quantities

	I	suspect	the	present	forecast	is	pessimistic	.
CD	3E-7							
DT			3E-8					
JJ		1E-9	1E-12	3E-12			7E-23	
NN	4E-6	2E-10	1E-13	6E-13	4e-16			
NNP	1E-5		4E-13					
NNS						1E-21		
PRP	4E-3							
RB				2E-14				
VB		6E-9		3E-15	2E-19			
VBD					6E-18			
VBN					4E-18			
VBP		5E-7	4E-14	4E-15	9E-19			
VBZ						6E-18		
.								2E-24

Decoding with the Viterbi Quantities

	I	suspect	the	present	forecast	is	pessimistic	.
CD	3E-7							
DT			3E-8					
JJ		1E-9	1E-12	3E-12			7E-23	
NN	4E-6	2E-10	1E-13	6E-13	4e-16			
NNP	1E-5		4E-13					
NNS						1E-21		
PRP	4E-3							
RB				2E-14				
VB		6E-9		3E-15	2E-19			
VBD					6E-18			
VBN					4E-18			
VBP		5E-7	4E-14	4E-15	9E-19			
VBZ						6E-18		
.								2E-24

Full Viterbi Algorithm

```
viterbi( $x, \Lambda, P_{\text{init}}, P_{\text{trans}}, P_{\text{final}}, P_{\text{emiss}}$ )  
  # Forward pass  
  for  $c_k \in \Lambda$   
    viterbi( $1, x, c_k$ ) =  $P_{\text{init}}(c_k | \text{start}) \cdot P_{\text{emiss}}(x_1 | c_k)$   
  for  $i = 2$  to  $N$ :  
    for  $c_k \in \Lambda$ :  
      viterbi( $i, x, c_k$ ) =  $\left( \max_{c_l \in \Lambda} P_{\text{trans}}(c_k | c_l) \cdot \text{viterbi}(i-1, x, c_l) \right) \cdot P_{\text{emiss}}(x_i | c_k)$   
      backtrack( $i, x, c_k$ ) =  $\left( \arg \max_{c_l \in \Lambda} P_{\text{trans}}(c_k | c_l) \cdot \text{viterbi}(i-1, x, c_l) \right)$   
  viterbi( $N+1, x, \text{stop}$ ) :=  $\max_{c_l \in \Lambda} P_{\text{final}}(\text{stop} | c_l) \cdot \text{viterbi}(N, x, c_l)$   
  backtrack( $N+1, x, \text{stop}$ ) =  $\arg \max_{c_l \in \Lambda} P_{\text{final}}(\text{stop} | c_l) \cdot \text{viterbi}(N, x, c_l)$   
  
  # Backward pass  
   $\hat{y}_N = \text{backtrack}(N+1, x, \text{stop})$   
  for  $i = N-1$  to  $1$ :  
     $\hat{y}_i = \text{backtrack}(i+1, x, \hat{y}_{i+1})$ 
```

Viterbi Summary

- Let N be the length of the sequence and m the number of states
- Memory:
 - The Viterbi table contains $N \times m$ numbers
 - The Backtrack table contains $N \times m$ numbers
- Runtime:
 - Each cell in the table requires $O(m)$ operations
 - Total runtime is $O(N \cdot m^2)$

Viterbi vs Posterior Decoding

- Is Viterbi decoding better than posterior decoding?
- Imagine the game: given a sequence x and an HMM
- G1) All or nothing Cost (Viterbi decoding)
 - Pay 10\$ if you make a single error
 - You care a lot about the whole label sequence coherency
- G2) Hamming Cost (Posterior decoding)
 - Pay 1\$ for every error in your decoding
 - You don't care much about the whole label sequence coherency